

EMPLOYEE ATTENDANCE USING FACIAL RECOGNITION

Abdulrahman Alrubaiya
Hatim Alshehri
Mohammed Alghamdi

Table of Content

- Introduction
- Objective
- Tools
- Models
- Lessons learned




Introduction

- Facial recognition is important in different domains
 - User experience
 - Public safety and national security
 - Control access
- Physiological biometric
 - Fingerprint
 - Iris etc..
- Can be used on existing hardware infrastructure
 - Cameras, image capturing devices



Objective


- Current pandemic
 - Contactless
- Employee attendance
 - Liveness / spoofed
 - Emotions detection
- Alerts
 - Unauthorized personnel
 - Motivational quotes
 - Mediums
 - SMS, Email etc..



TOOLS




Models:



01

Model: Liveness


- Objective
- Data
- Baseline
- Network architecture
- Eval metrics
- Challenges



02

Model: Facial Recognition

- Objective
- Data
- Baseline
- Network architecture
- Eval metrics
- Challenges




03

Model: Emotion detection

- Objective
- Data
- Baseline
- Network architecture
- Eval metrics
- Challenges


Liveness detection

- Objective: Create liveness detector capable of identifying real and fake/spoofed faces
- Data:
 - 1 minute video on iPhone for real, 1 minute video on Logitech Webcam for fake
 - 3600 images(1800 for each class), 300x300



Real (1080x1920)

Fake/Spoofed (960x720)



Liveness (Cont.)

- Preprocessing:
- Resize image to 32x32
- Intensify images to range of 0, 1
- Label encode (real, fake)
- One hot encode
- Split data train, validation 75-25

Data augmentation

- Rotation 20
- Zoom 0.15
- Width shift 0.2
- Height shift 0.2
- Horizontal flip




Liveness (Cont.)


- Baseline:
 - Simple CNN
 - 32, 16, D(24), 2 →
Kernel=3, ReLU,
Maxpooling 2x2,
softmax

Set	Score
Training	0.86
Validation	0.84

Real



Fake



Liveness using LivenessNet (Cont.)

- LivenessNet

- Hidden layers: 6
- First:
 - Conv2D_1: Filters=16, kernel_size=3, padding=same
 - Activation=ReLU
 - BatchNormalization
 - MaxPooling=2x2
 - Dropout=0.25
- Second:
 - Conv2D_2: Filters=32, kernel_size=3, padding=same
 - Activation=ReLU
 - BatchNormalization
 - MaxPooling=2x2
 - Dropout=0.25

First: Conv => ReLU => Conv => ReLU=> Pool

Second: Conv => ReLU=> Conv => ReLU => Pool

Flatten set: Flatten() => Dense(64) => ReLU

Softmax classifier: Dense(2) => Softmax

- Loss=binary_crossentropy
- Optimizer=adam
- Metrics=accuracy
- Epochs=50, Batch_size=8

Liveness using LivenessNet (Cont.)

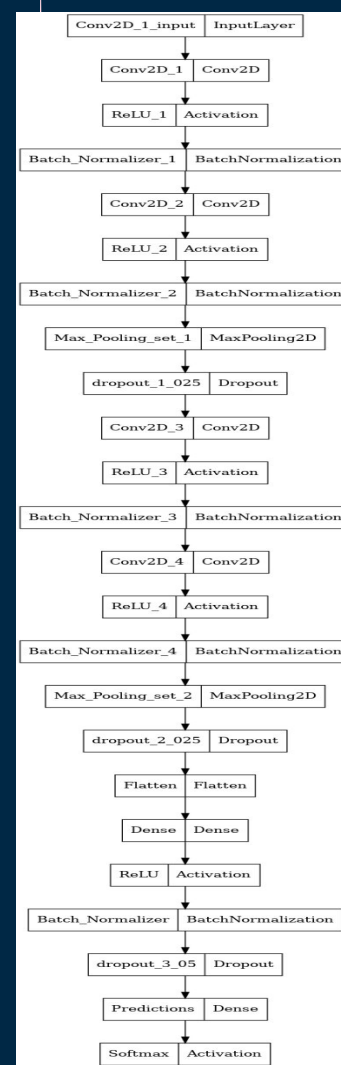
```
In [3]: # model summary  
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
Conv2D_1 (Conv2D)	(None, 32, 32, 16)	448
ReLU_1 (Activation)	(None, 32, 32, 16)	0
Batch Normalizer_1 (Batch Normalization)	(None, 32, 32, 16)	64
Conv2D_2 (Conv2D)	(None, 32, 32, 16)	2320
ReLU_2 (Activation)	(None, 32, 32, 16)	0
Batch Normalizer_2 (Batch Normalization)	(None, 32, 32, 16)	64
Max Pooling_set_1 (MaxPooling2D)	(None, 16, 16, 16)	0
dropout_1_025 (Dropout)	(None, 16, 16, 16)	0
Conv2D_3 (Conv2D)	(None, 16, 16, 32)	4640
ReLU_3 (Activation)	(None, 16, 16, 32)	0
Batch Normalizer_3 (Batch Normalization)	(None, 16, 16, 32)	128
Conv2D_4 (Conv2D)	(None, 16, 16, 32)	9248
ReLU_4 (Activation)	(None, 16, 16, 32)	0
Batch Normalizer_4 (Batch Normalization)	(None, 16, 16, 32)	128
Max Pooling_set_2 (MaxPooling2D)	(None, 8, 8, 32)	0
dropout_2_025 (Dropout)	(None, 8, 8, 32)	0
Flatten (Flatten)	(None, 2048)	0
Dense (Dense)	(None, 64)	131136
ReLU (Activation)	(None, 64)	0
Batch Normalizer (Batch Normalization)	(None, 64)	256
dropout_3_05 (Dropout)	(None, 64)	0
Predictions (Dense)	(None, 2)	130
Softmax (Activation)	(None, 2)	0


```
=====
Total params: 148,562  
Trainable params: 148,242  
Non-trainable params: 320
```

Liveness using LivenessNet (Cont.)



Liveness using LivenessNet (Cont.)

Set	Score
Training	1
Validation	1



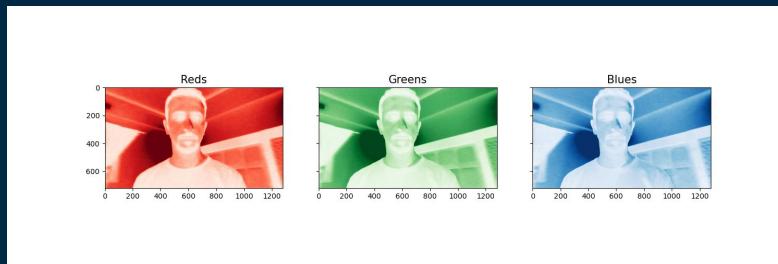
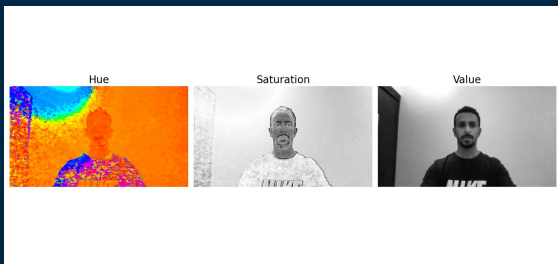
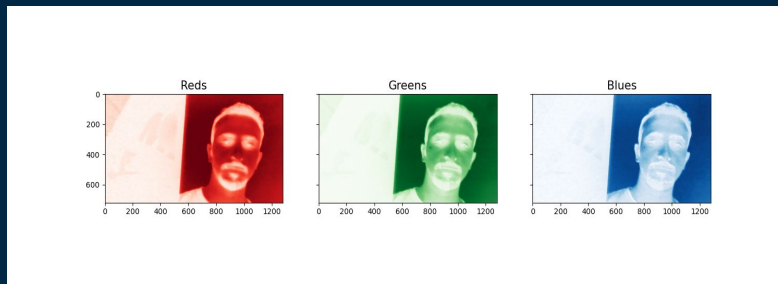

Liveness challenges

- Unable to generalize → Took frame by frame pictures instead of 1 frame per second
- White-lit room provides false predictions → Try different models



How problem solved

No good result with these images





3400 Real

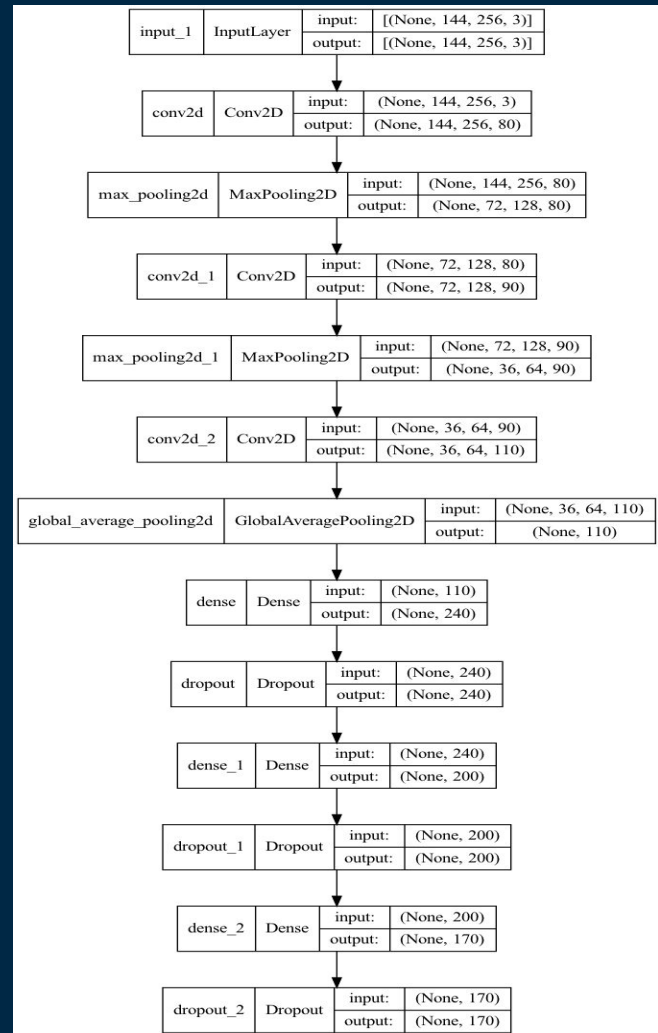
3400 with phone

How problem solved

Using “Fast two-dimensional phase- unwrapping algorithm based on sorting by reliability following a noncontinuous path , 2002“ paper



- accuracy: 0.9952
- val_accuracy: 0.9793
- Epoch 00057: early stopping
- Test accuracy: 0.9828






Facial Recognition Model:

DATASET

- Facial Recognition.
- From OpenCv (ImageCapture).
- 9,000
- Multilabel:
 - Hatim
 - Abdulrahman
 - Mohammed

LABEL



DATA AUGMENTATION

- Resize Images (220X220)
- Scaling/Normalization (1 to 255)
- Randomly Rotate Images (up to 20°)
- Randomly Flip Images (Horizontally)

Train



80%

Validation



20%

DATA SAMPLE



Facial Recognition Models:



Simple NN

- Hidden layers: 1
- Dense: 16
- Activation: ReLU
- Optimizer: ADAM
- Loss function: Categorical Crossentropy
- Output Activation: SoftMax



CNN

- Layers added : 4
- Dense : 1000 – 850 – 1000 – 500
- BatchNormalization
- Activation : ReLU
- Optimizer : adam
- Weights : imagenet
- Include Top : False
- Loss function: Categorical Crossentropy
- Output Activation: SoftMax



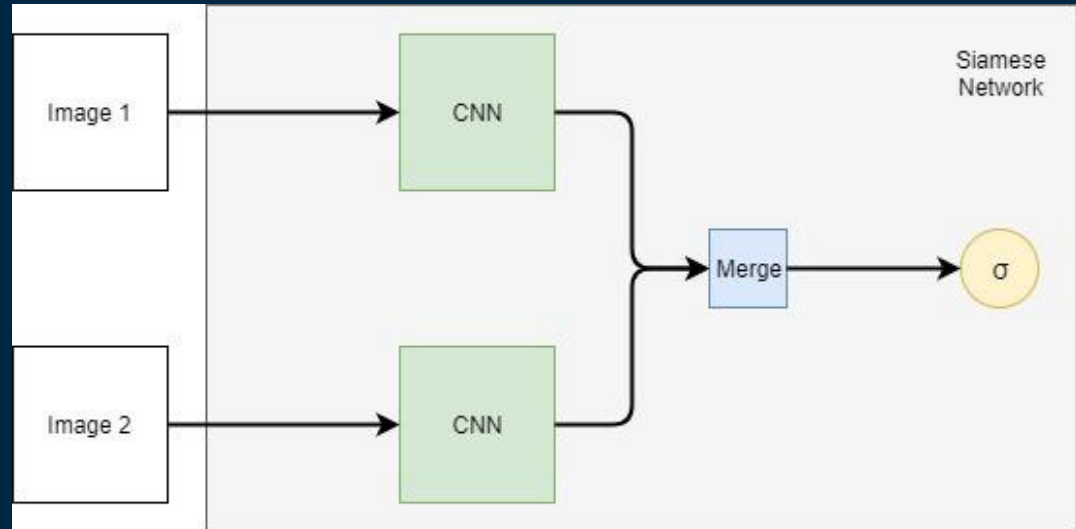
Scores

- Baseline Simple NN:
 - Train Accuracy: 100%
 - Validation Accuracy: 84.76%
-
- CNN
 - Train Accuracy: 100%
 - Validation Accuracy: 84.76%

Siamese Neural Networks Model:

After we tried the previous model in order to perform the facial recognition and got unsatisfactory results we decided to choose another approach to make this mission work well.

Siamese NN Architecture:



Data

- Bast Data : 17432 images
- 105 class
- Used :
- 23469 Images in total
- 4 classes



Model : Siamese Network

- Training accuracy: 0.9987
- val_accuracy: 0.99875
- Test accuracy: 0.9987

CNN


- Conv2D : 3
- Filter: 50 - 80 - 90
- Layers added : 3
- Dense : 150 - 190 - 120
- BatchNormalization
- Activation : ReLU
- Optimizer : adam

Similarity Example

.99



.018



The background is a dark blue gradient. It features several thin, vertical white lines of varying lengths. Scattered across the background are numerous small squares in three colors: light blue, light orange, and light pink. Some of these squares are solid, while others are outlined. The overall aesthetic is modern and minimalist.

Emotion Detection Model

DATASET

- Facial Emotion Recognition
- From Kaggle
- ~ 37K
- Multilabel:
 - Neutral
 - Happy
 - Sad
 - Surprise
 - Fear
 - Disgust
 - Anger
- Greyscale



DATA PRE-PROCESSING

- Resize Images (142X256)
- Scaling/Normalization (0,1)
- Randomly Rotate Images (up to 20°)
- Randomly Flip Images (Horizontally)

Train




80%

Validation



20%

DATA SAMPLE



Transfer Learning MODELS:



MobileNet

- Layers added : 4
- Dense : 1000 – 850 – 1000 - 500
- BatchNormalization
- Dropout : 0.3
- Activation : ReLU .
- Optimizer : adam

- Loss function: Categorical Crossentropy
- Output Activation: SoftMax

- Train Accuracy: 86.29%
- Validation Accuracy: 48.71%



VGG-16

- Layers added : 3
- Dense : 1000 – 256 - 128
- BatchNormalization
- Dropout : 0.5
- Activation : ReLU.
- Optimizer : adam

- Train Accuracy: 83.46%
- Validation Accuracy: 57.92%

Lessons learned

- Version control branch merging
- Github issues to track progress
- Academic papers
 - Stanford method
 - Three pass
 - 1: Title, abstract, introduction, headings, sub-headings, and conclusion
 - 2: Read and look at figures carefully
 - 3: Read paper carefully with all details

The background is a dark navy blue. It is decorated with various geometric elements: small squares in teal, orange, and pink, and thin white vertical lines of varying lengths. These elements are scattered across the frame, creating a modern, minimalist aesthetic.

THANK YOU
FOR
LISTENING